

## TABLE OF CONTENTS

1. Summary .....	3
2. Solution and Implementation Details .....	4
a. Setting Up the Environment .....	4
b. Code .....	12
3. Results .....	24
a. Kibana .....	24
b. Cassandra .....	24
4. Conclusion and Lessons Learned .....	27

**Literature is included as end notes to this paper,  
referring to the respective paragraphs.**

# 1. Summary

## Project Goal and Problem Statement

This project's goal is to develop an integrated framework that analyzes data from an IOT smartwatch. It will provide a platform to gain insights into raw data as well as aggregated metrics.

## Big Data Source

Kafka will be used as data generator to produce a constant message stream simulating incoming data from an IOT smartwatch/fitness tracker.

## Results

As a result of my employed architecture, I enabled the user to show and analyze (a) raw data features (heart rate, stress level) as a function of time and (b) aggregated metrics (e.g. total number of users aggregated by different sports in total and within the last 10 minutes).

## Processing Pipeline



## Pipeline Overview and Technologies Used

**Collection Tier:** Apache Kafka Producer will be used to simulate a stream of incoming messages by multiple fitness tracker devices used in Frankfurt.

**Processing Tier:** Spark Streaming will be used to calculate total number of users and the total number of users aggregated by different sports within the last hour.

**Storage Tier:** Apache Cassandra will be used to store the calculated metrics as tables in a database.

**Indexing Tier:** Kafka ES Connector will be used to push data from the initial Kafka simulator to Elasticsearch.

**Search Tier/Storage Tier:** Elasticsearch will be used to act as a storage for the raw data. An event will be the signal send by a smartwatch, where each event will receive a unique index for logging purposes.

**Visualization Tier:** Kibana will be used to produce time-series graphs showing the data for heart rate and stress level over time for different users

## 2. Implementation Details

### a. Setting Up the Environment

I employed the architecture on Windows 10, but it should – with minor adjustments – work on other operating systems as well. This architecture requires the following software components and versions:

Software/Packages	Version
Java Development Kit	1.8
Maven	3.5.2
Eclipse	Java Oxygen
Zookeeper	3.4.9
Kafka	2.10 // 0.10.0.0
Cassandra	3.11.1
Spark	1.6.2
Elasticsearch	5.0.0
Kibana	5.0.0
Kafka ES Connector	4.0.0

I tried to previously implement the solution in Python but ended up abandoning this approach and switched to Java because the **Python support** for an integration of Spark Streaming and Kafka was **discontinued**. Please refer to <https://spark.apache.org/docs/2.2.0/streaming-kafka-integration.html> for more information.

Furthermore, I ran into **issues** with integrating the **Kafka ES Connector** with **Elasticsearch 6.0.0** which is why I switched to Elasticsearch 5.0.0. The reason for this is that Elasticsearch 6.0.0 deprecated the use of multiple mapping types which are needed to use the connector. However, they are still supported in version 5.0.0, so I just used that instead of reprogramming the connector. Please refer to <https://www.elastic.co/guide/en/elasticsearch/reference/6.x/removal-of-types.html> for more information.

With regards to the other software components, please follow the installation guidelines given below to install them on your respective operating system:

JDK:

[https://docs.oracle.com/javase/8/docs/technotes/guides/install/install\\_overview.html](https://docs.oracle.com/javase/8/docs/technotes/guides/install/install_overview.html)

Maven:

<https://maven.apache.org/install.html>

Eclipse:

[http://www.eclipse.org/downloads/eclipse-packages/?show\\_instructions=TRUE](http://www.eclipse.org/downloads/eclipse-packages/?show_instructions=TRUE)

Zookeeper:

<https://zookeeper.apache.org/doc/r3.1.2/zookeeperStarted.html>

Kafka:

<https://kafka.apache.org/quickstart>

Cassandra:

[http://cassandra.apache.org/doc/latest/getting\\_started/installing.html](http://cassandra.apache.org/doc/latest/getting_started/installing.html)

Spark:

<https://spark.apache.org/docs/1.6.2>

Elasticsearch:

<https://www.elastic.co/guide/en/elasticsearch/reference/5.0/install-elasticsearch.html>

Kibana:

<https://www.elastic.co/guide/en/kibana/5.0/install.html>

Kafka ES Connector:

<https://github.com/confluentinc/kafka-connect-elasticsearch>

In the following, I describe the process to get Zookeeper, Apache Kafka, Apache Cassandra, Elasticsearch, Kibana and the Kafka ES Connector running. I will not go into further details with regards to JDK, Maven, Eclipse and Spark. Please refer to the guidelines for help on installation issues.

## Zookeeper

After downloading and unzipping the Zookeeper package, the configuration file “zoo\_sample.cfg” needs to be renamed to “zoo.cfg” and

```
dataDir=/tmp/zookeeper
```

needs to be changed to

```
dataDir=\\Zookeeper\\data
```

Furthermore, the following entries should be added to the system environment variables:

```
ZOOKEEPER_HOME = C:\\Zookeeper
```

and

```
PATH = %ZOOKEEPER_HOME%\\bin
```

Eventually, the Zookeeper server can be started by typing

```
> zkserver
```

in the Windows command prompt:

## Command Prompt

```
Microsoft Windows [Version 10.0.15063]
(c) 2017 Microsoft Corporation. All rights reserved.

C:\Users\akreu_000>cd C:\Zookeeper

C:\Zookeeper>zkservice
```

```
2017-12-11 13:02:56,524 [myid:] - INFO [main:Environment@100] - Server environm
nt:user.name=akreu_000
2017-12-11 13:02:56,527 [myid:] - INFO [main:Environment@100] - Server environm
nt:user.home=C:\Users\akreu_000
2017-12-11 13:02:56,590 [myid:] - INFO [main:ZooKeeperServer@825] - tickTime se
t to 2000
2017-12-11 13:02:56,595 [myid:] - INFO [main:ZooKeeperServer@834] - minSessionT
imeout set to -1
2017-12-11 13:02:56,596 [myid:] - INFO [main:ZooKeeperServer@843] - maxSessionT
imeout set to -1
2017-12-11 13:02:57,648 [myid:] - INFO [main:ServerCnxnFactory@117] - Using org
.apache.zookeeper.server.NIOServerCnxnFactory as server connection factory
2017-12-11 13:02:57,663 [myid:] - INFO [main:NIOServerCnxnFactory@89] - binding
to port 0.0.0.0/0.0.0.0:2181
```

## Kafka

After downloading and unzipping the Kafka package, the 'server.properties' file in the config folder. Again, change

```
> log.dirs=/tmp/kafka-logs
```

to

```
> log.dir=C:\Kafka\kafka-logs
```

Kafka will run on the default port 9092 and Zookeeper will run on 2181. Afterwards, the Kafka server can be started by invoking

```
> cd C:\Kafka
> .\bin\windows\kafka-server-start.bat .\config\server.properties
```

in the Windows command line or

```
> .\bin\windows\kafka-server-start.sh .\config\server.properties
```

in the Cygwin tool.

The console output look like the following to indicate a successful start of the Kafka server:

```
[2017-12-11 13:39:35,829] INFO [Partition __consumer_offsets-32 broker=0] __cons
umer_offsets-32 starts at Leader Epoch 3 from offset 0. Previous Leader Epoch wa
s: 2 (kafka.cluster.Partition)
[2017-12-11 13:39:35,830] INFO [Partition __consumer_offsets-3 broker=0] __consu
mer_offsets-3 starts at Leader Epoch 3 from offset 0. Previous Leader Epoch was:
2 (kafka.cluster.Partition)
[2017-12-11 13:39:35,831] INFO [Partition __consumer_offsets-13 broker=0] __cons
umer_offsets-13 starts at Leader Epoch 3 from offset 0. Previous Leader Epoch wa
s: 2 (kafka.cluster.Partition)
[2017-12-11 13:49:24,923] INFO [GroupMetadataManager brokerId=0] Removed 0 expir
ed offsets in 0 milliseconds. (kafka.coordinator.group.GroupMetadataManager)
```

In a second command prompt, create a topic 'iotdata' by invoking the following command:

```
> .\bin\windows\kafka-topics.bat --create --zookeeper localhost:2181 --replication-factor 1 --partitions 1 --topic iotdata
```

## Cassandra

After downloading the most recent Cassandra version from [Cassandra.apache.org/download](https://cassandra.apache.org/download/), unzip the package and move it to your desired Cassandra directory.



## Downloading Cassandra

### Latest version

Download the latest Apache Cassandra 3.11 release: [3.11.1](#) ([pgp](#), [md5](#) and [sha1](#)), released on 2017-10-10.

To start Cassandra, invoke the following commands in the command prompt:

```
> cd C:\Cassandra\bin
> .\cassandra -f
```

Navigate to a new terminal and execute

```
> cqlsh
```

to start the command line tool to insert CQL statements and interact with your Cassandra database.

We need to create a key space and two different<sup>i</sup> tables<sup>ii</sup> before starting to execute our smartwatch simulation. We do this by invoking the following commands in the CQLSH prompt<sup>iii</sup>:

```
> CREATE KEYSPACE IF NOT EXISTS IOTKeySpace WITH
  replication = {'class':'SimpleStrategy',
    'replication_factor':1};
> CREATE TABLE IOTKeySpace.TotalUsersBySports (sports
```

```

text, count int, timestamp
timestamp, date text, PRIMARY KEY
(sports, date));

> CREATE TABLE IOTKeySpace.UsersBySports_Last10Seconds
(sports
text, count int, timestamp
timestamp, date text, PRIMARY KEY
(sports, date));

```

The primary keys were chosen to support the inherent structure of the queries and optimize queuing time.

```

cqlsh> CREATE KEYSPACE IF NOT EXISTS IOTKeySpace WITH replication = {'class':'SimpleStrategy', 'replication_factor':1};
cqlsh> CREATE TABLE IOTKeySpace.TotalUsersBySports (sports text, count int, timestamp timestamp, date text, PRIMARY KEY (sports, date));
cqlsh> CREATE TABLE IOTKeySpace.UsersBySports_Last10Seconds (sports text, count int, timestamp timestamp, date text, PRIMARY KEY (sports, date));
cqlsh>

```

## Elasticsearch

The setup of Elasticsearch is straight-forward. After downloading and unzipping the respective package for Elasticsearch 5.0.0 (IMPORTANT! Don't use 6.0.0, see above for the reasoning behind this), the following commands can be used to start Elasticsearch:

```

> cd C:\Elasticsearch
> .\bin\elasticsearch

```

The console output should look like this:

```

[2017-12-12T08:53:47,668][INFO ][o.e.c.m.MetadataCreateIndexService] [1MYXH-8] [iotdata] creating index, cause [api], templates [], shards [5]/[1], mappings [iotdata]
[2017-12-12T08:55:05,728][INFO ][o.e.c.m.MetadataMappingService] [1MYXH-8] [iotdata/SZT0-r4jQqeE3wy3fROvwQ] create_mapping [log]
[2017-12-12T08:55:49,365][INFO ][o.e.c.m.MetadataMappingService] [1MYXH-8] [.kibana/vbSaFSA4SIWWCtkUIs-a_w] create_mapping [dashboard]
[2017-12-12T08:56:01,330][INFO ][o.e.c.m.MetadataMappingService] [1MYXH-8] [.kibana/vbSaFSA4SIWWCtkUIs-a_w] create_mapping [timelion-sheet]
[2017-12-12T09:12:02,431][INFO ][o.e.c.m.MetadataCreateIndexService] [1MYXH-8] [iotdata] creating index, cause [api], templates [], shards [5]/[1], mappings [iotdata]
[2017-12-12T09:12:17,380][INFO ][o.e.c.m.MetadataMappingService] [1MYXH-8] [iotdata/MjAOCQbbQ6yFRrkCG6J76Q] create_mapping [log]
[2017-12-12T09:20:11,527][INFO ][o.e.c.m.MetadataMappingService] [1MYXH-8] [.kibana/vbSaFSA4SIWWCtkUIs-a_w] create_mapping [visualization]

```

Use a simple PUT request to ensure that Elasticsearch is up and running:

```

curl -XPUT 'http://localhost:9200/_search

```

## Kibana

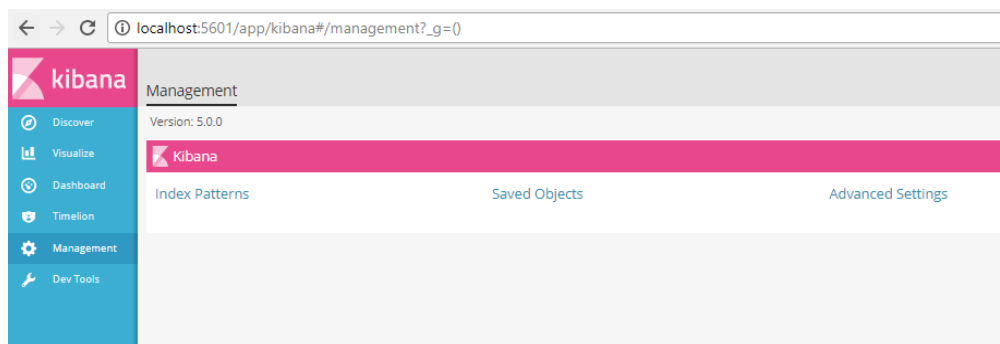
As simple to setup as Elasticsearch. However, it is important to start the Elasticsearch server before starting Kibana. Also, we need to use Kibana 5.0.0 instead of 6.0.0 to ensure compatibility within the architecture. After downloading and unzipping the respective package for Kibana 5.0.0, the following commands can be used to start Kibana:

```
> cd C:\Elasticsearch\Kibana
> .\bin\kibana
```

The console output should look like this:

```
C:\kibana-5.0.0-windows-x86>.\bin\kibana
log [06:40:29.967] [info][status][plugin:kibana@5.0.0] Status changed from
nitialized to green - Ready
log [06:40:30.165] [info][status][plugin:elasticsearch@5.0.0] Status change
from uninitialized to yellow - Waiting for Elasticsearch
log [06:40:30.276] [info][status][plugin:console@5.0.0] Status changed from
uninitialized to green - Ready
log [06:40:32.488] [info][status][plugin:timelion@5.0.0] Status changed fro
uninitialized to green - Ready
log [06:40:32.526] [info][listening] Server running at http://localhost:560
log [06:40:32.535] [info][status][ui settings] Status changed from uninitia
ized to yellow - Elasticsearch plugin is yellow
log [06:40:37.627] [info][status][plugin:elasticsearch@5.0.0] Status change
from yellow to yellow - No existing Kibana index found
log [06:40:39.965] [info][status][plugin:elasticsearch@5.0.0] Status change
from yellow to green - Kibana index ready
log [06:40:39.970] [info][status][ui settings] Status changed from yellow t
green - Ready
```

If everything worked according to plan, Kibana can be reached under <http://localhost:5601> from a browser of your choice. You should see a working Kibana environment:



Once Kibana is active, we use the included Dev Tools (left menu, last entry) to create an index `iotdata_index` and assign it a mapping, so that Elasticsearch knows right away which data types are being sent via the stream:



The mapping is:

```
PUT iotdata
{
  "mappings": {
```



```

    "iotdata": {
      "properties": {
        "userid": { "type": "keyword"},
        "sports": { "type": "keyword"},
        "heartrate": { "type": "integer"},
        "stresslevel": { "type": "integer"},
        "timestamp": {
          "type": "date",
          "format": "yyyy-MM-dd HH:mm:ss"
        }
      }
    }
  }
}

```

## Kafka Elasticsearch Connector

Starting the Kafka Elasticsearch Connector is a bit more involved process<sup>iv</sup>. First, the version compatible with Kafka version 0.10.0.0 needs to be cloned from github using

```

> git clone -b 0.10.0.0
https://github.com/confluentinc/kafka-connect-
elasticsearch.git

```

```

C:\Users\akreu_000>git clone -b 0.10.0.0 https://github.com/confluentinc/kafka-c
onnect-elasticsearch.git
Cloning into 'kafka-connect-elasticsearch'...
remote: Counting objects: 1854, done.
remote: Compressing objects: 100% (50/50), done.
remote: Total 1854 (delta 23), reused 58 (delta 12), pack-reused 1776
Receiving objects: 100% (1854/1854), 329.31 KiB | 248.00 KiB/s, done.
Resolving deltas: 100% (762/762), done.
C:\Users\akreu_000>

```

Afterwards, we build the package by invoking the following commands<sup>v</sup> in the command line:

```

> cd kafka-connect-elasticsearch
> mvn clean package

```

This should lead to the following output:

```

[INFO] Reading assembly descriptor: src/assembly/development.xml
[INFO] Reading assembly descriptor: src/assembly/package.xml
[INFO] Copying files to C:\kafka-connect-elasticsearch\target\kafka-connect-e
lasticsearch-3.2.0-SNAPSHOT-development
[WARNING] Assembly file: C:\kafka-connect-elasticsearch\target\kafka-connect-
elasticsearch-3.2.0-SNAPSHOT-development is not a regular file (it may be a dire
ctory). It cannot be attached to the project build for installation or deployment.
[INFO] Copying files to C:\kafka-connect-elasticsearch\target\kafka-connect-e
lasticsearch-3.2.0-SNAPSHOT-package
[WARNING] Assembly file: C:\kafka-connect-elasticsearch\target\kafka-connect-
elasticsearch-3.2.0-SNAPSHOT-package is not a regular file (it may be a director
y). It cannot be attached to the project build for installation or deployment.
[INFO] -----
[INFO] BUILD SUCCESS
[INFO] -----
[INFO] Total time: 01:34 min
[INFO] Finished at: 2017-12-11T22:33:19-05:00
[INFO] Final Memory: 36M/418M
[INFO] -----

```

After the build is finished, go to the *target/kafka-connect-elasticsearch- 3.2.0-SNAPSHOT-package/share/java/kafka-connect-elasticsearch/* directory and copy all the files in there into the Kafka *libs* directory.

We will run the ES connector in standalone mode. To be able to do that, we need to create a new configuration file, *elasticsearch-connect.properties* in the *conf*-folder of the Kafka directory. In that file, insert

```
name=elasticsearch-sink
connector.class=io.confluent.connect.elasticsearch.ElasticsearchSinkConnector
tasks.max=1
topics=iotdata
topic.index.map=iotdata:iotdata_index
connection.url=http://localhost:9200
type.name=log
key.ignore=true
schema.ignore=true
```

*key.ignore* and *schema.ignore* are set to *true* because we already assigned a mapping in *elasticsearch* to that index.

Furthermore, we need to alter the *connect-standalone.properties* file in the Kafka *conf* directory. It should include the following content:

```
bootstrap.servers=localhost:9092
key.converter=org.apache.kafka.connect.json.JsonConverter
value.converter=org.apache.kafka.connect.json.JsonConverter
key.converter.schemas.enable=false
value.converter.schemas.enable=false
internal.key.converter=org.apache.kafka.connect.json.JsonConverter
internal.value.converter=org.apache.kafka.connect.json.JsonConverter
internal.key.converter.schemas.enable=false
internal.value.converter.schemas.enable=false
offset.storage.file.filename=/tmp/connect.offsets
offset.flush.interval.ms=10000
```

An important thing to notice here is that the way we set up the connector, we can only parse JSON streams with it, no other format.

We start the connector with the following command:

```
> bin\connect-standalone.sh config\connect-standalone.properties config\elasticsearch-connect.properties
```

Elasticsearch will now “subscribe” to topic “iotdata” and every message that gets sent to this topic will be directed to and parsed by Elasticsearch.

## **b. Code**

In the following, the JAVA classes<sup>vi</sup> necessary to simulate the smartwatch data<sup>vii</sup> and to send it to the 'iotdata' topic are explained.

### **Package: watchdata.producer**

This package is responsible for creating Smartwatch data and sending it via the Kafka Producer.

#### **Smartwatch class**

The Smartwatch class is the cornerstone of the data producer. It is responsible for connecting to Kafka and Zookeeper and it creates events in the required JSON format before sending them off via the Kafka Producer to the 'iotdata' topic.

#### **Data Encoder**

The Data Encoder Class encodes the events.

### **Package: watchdata.processor**

The processor package is used to parse and process the data from the Kafka Producer. It is responsible for calculating aggregated metrics and sending them to the Cassandra database.

#### **TotalUserData**

In this class, the object structure for the first of the two aggregated measures is set: The count of total users per sports.

#### **WindowUserData**

In this class, the object structure for the second of the two aggregated measures is set: The count of total users per sports within the last 30 seconds.

#### **DataProcessor**

The data processor is the core of the processor package. Here, we will parse the incoming JSON objects and calculate our metrics before writing the results to Cassandra. For the first metric, the total user count by sports parsed by date, we need to process both current streaming data and historical data of that day. For the second measure, which counts the users per sports within the last 30 seconds, we need to employ a windowed datastream.

#### **DataDecoder**

The Data Decoder class corresponds to the Data Encoder Class within the producer package and is responsible for decoding the incoming objects.

## **Data**

This class is a copy of the Data Class within the producer package. Again, it gives structure to the data that the smartwatches are sending.

## **Aggregate Key**

This class is used to construct aggregate for mapping the data streams in Cassandra.

## **Data Streamer**

This class sets the Spark-Context and parses data from the incoming JSON stream of events sent by the smartwatches. It furthermore ensures that the incoming stream batches are being reduced by key (here: userId) to allow for counting unique users in the aggregated metrics we'll later store in Cassandra.

### 3. Results

The code works as expected. In Cassandra, we start the CQLSH console by executing the following command in the Windows CLI:

```
> cd \path\to\Cassandra\directory\bin
> cqlsh
> SELECT * FROM iotkeyspace.totalusersbysports;
```

```
cqlsh> SELECT * FROM iotkeyspace.totalusersbysports;

sports | date       | count | timestamp
-----+-----+-----+-----
skiing | 2017-12-12 | 5      | 2017-12-12 21:27:00.384000+0000
biking | 2017-12-12 | 7      | 2017-12-12 21:25:58.053000+0000
swimming | 2017-12-12 | 4      | 2017-12-12 21:29:18.473000+0000
hiking | 2017-12-12 | 4      | 2017-12-12 21:28:48.689000+0000
running | 2017-12-12 | 2      | 2017-12-12 21:30:18.607000+0000

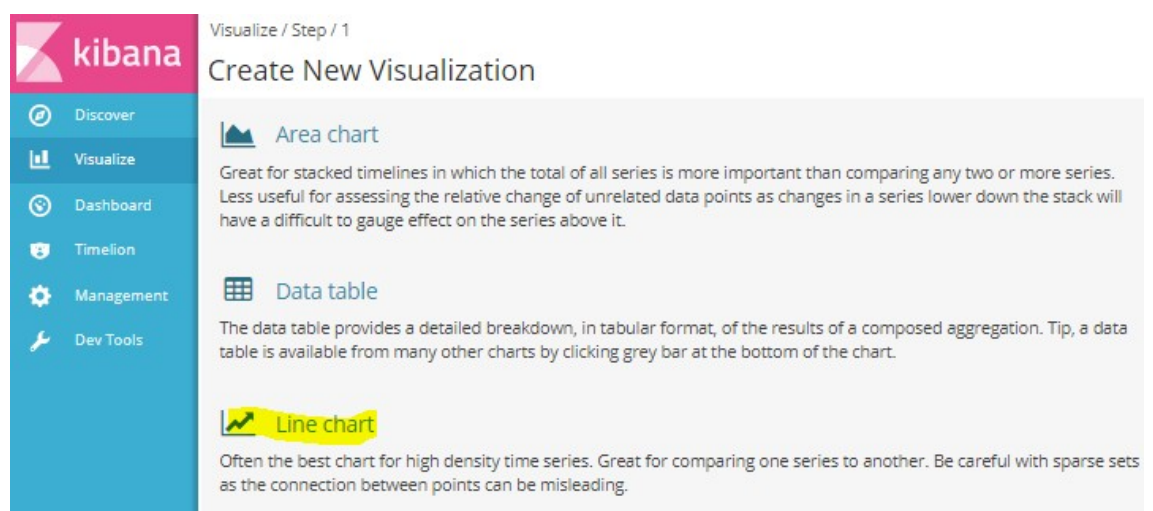
(5 rows)
```

```
> SELECT * FROM iotkeyspace.usersbysports_last10seconds;
```

```
cqlsh> SELECT * FROM iotkeyspace.usersbysports_last10seconds;

sports | date       | count | timestamp
-----+-----+-----+-----
skiing | 2017-12-12 | 1      | 2017-12-12 21:27:24.878000+0000
biking | 2017-12-12 | 1      | 2017-12-12 21:26:18.885000+0000
swimming | 2017-12-12 | 1      | 2017-12-12 21:29:42.515000+0000
hiking | 2017-12-12 | 1      | 2017-12-12 21:29:13.065000+0000
running | 2017-12-12 | 1      | 2017-12-12 21:30:46.301000+0000
```

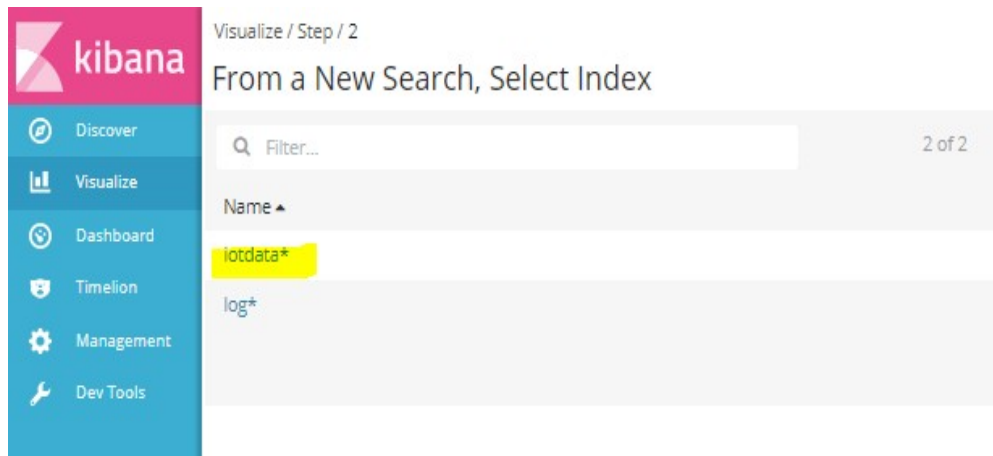
In Kibana, we navigate to “Visualize” on the left-hand side menu of the screen. We are prompted to create a new visualization. To obtain the time-series graphs we are interested in, we click on “Line Chart” and select our “iotdata” index.



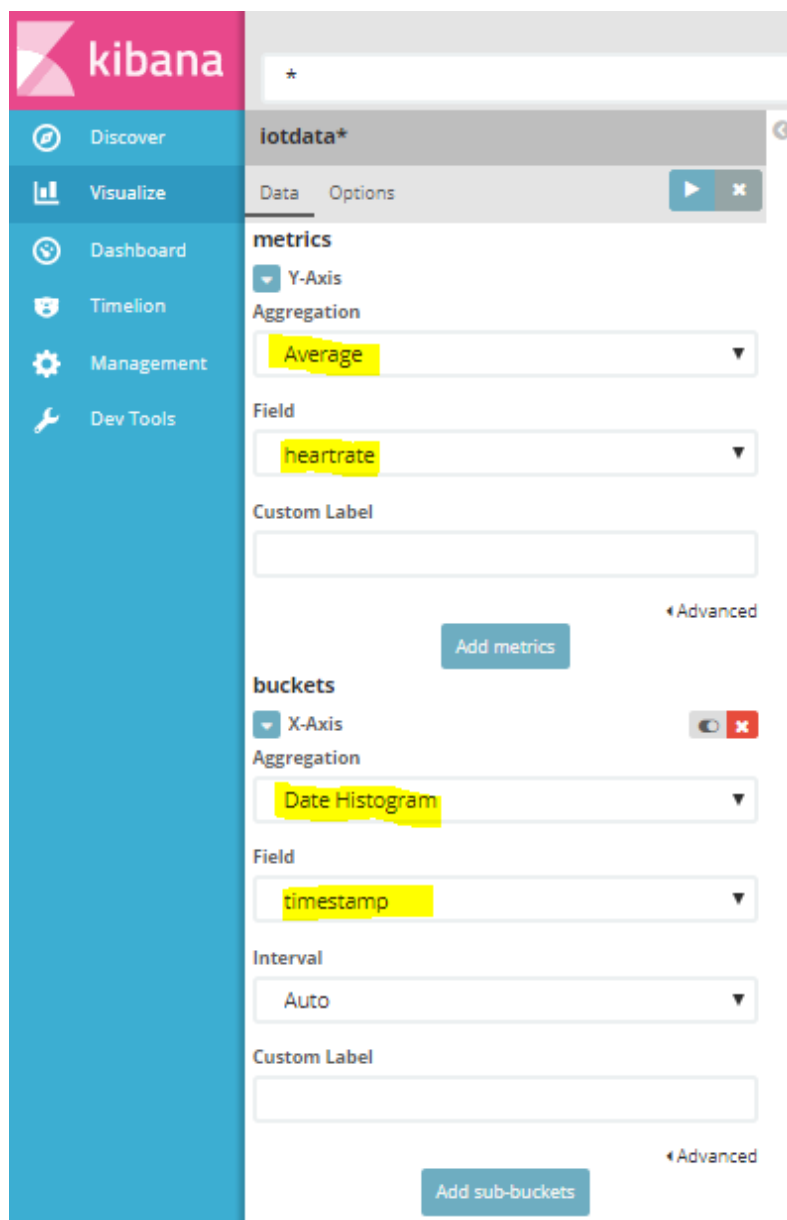
Visualize / Step / 1

## Create New Visualization

- Area chart**  
Great for stacked timelines in which the total of all series is more important than comparing any two or more series. Less useful for assessing the relative change of unrelated data points as changes in a series lower down the stack will have a difficult to gauge effect on the series above it.
- Data table**  
The data table provides a detailed breakdown, in tabular format, of the results of a composed aggregation. Tip, a data table is available from many other charts by clicking grey bar at the bottom of the chart.
- Line chart**  
Often the best chart for high density time series. Great for comparing one series to another. Be careful with sparse sets as the connection between points can be misleading.




Afterwards, we select “avg” and “heartrate” (or “stresslevel”, respectively) for the y-axis and “Date Histogram” and “timestamp” for the x-axis.




This yields to the following graphs:

Average of variable “heartrate” over time:



Average of variable “stresslevel” over time:



We successfully developed an architecture to visualize and analyte data from multiple smartwatches.

## 4. Conclusions and Lessons Learned

Originally, I wanted to include Re:dash as a visualization tool for Apache Cassandra. However, I missed that Re:dash doesn't run on my local machine but on remote servers. This means that I would've had to allow access from the internet to my database and to configure my data source accordingly. While this would have been possible, especially with a tool like Ngrok (<https://ngrok.com/>) to facilitate this, it would have meant an increased security risk. As I am no professional in security concerns, I did not feel comfortable to implement this solution and share it with the community, in case someone wanted to use it for professional purposes.

Another thing that I learned is that when planning an architecture, one not only has to pay attention to the high-level structure, i.e. how the individual components could fit together, but also to the version compatibility between the components. This cost me a lot of time and, as mentioned above, I started out with writing the code in Python until I realized that the Spark Streaming and Kafka integration support was discontinued, which let me start my project again in Java. So next time, I'll check version compatibility first before head starting into programming my code.

---

i <https://databricks.com/blog/2017/02/23/working-complex-data-formats-structured-streaming-apache-spark-2-1.html>

ii <http://abiasforaction.net/a-practical-introduction-to-cassandra-query-language/>

iii <https://www.slideshare.net/mattdennis/cassandra-nyc-2011-data-modeling>

iv <https://sematext.com/blog/kafka-connect-elasticsearch-how-to/>

v <https://docs.confluent.io/current/connect/userguide.html>

vi <https://acadgild.com/blog/streaming-twitter-data-using-kafka/>

vii <https://www.infoq.com/articles/traffic-data-monitoring-iot-kafka-and-spark-streaming>